

On-device AI - a short brief

Running large language models locally on consumer hardware is now practical. The benefits are privacy (your data never leaves the machine), zero cloud cost, and offline availability.

The QVAC SDK runs LLMs, embeddings, and LoRA fine-tuning directly on a Mac, iPhone or Android device. For Second Self this means your notes, your model, and your memory all stay on your computer.

Key tradeoff: smaller models are faster but less capable; pick the largest model your hardware can run comfortably.